



LangChain — Chat with your data

course link: <https://www.deeplearning.ai/short-courses/langchain-chat-with-your-data/>

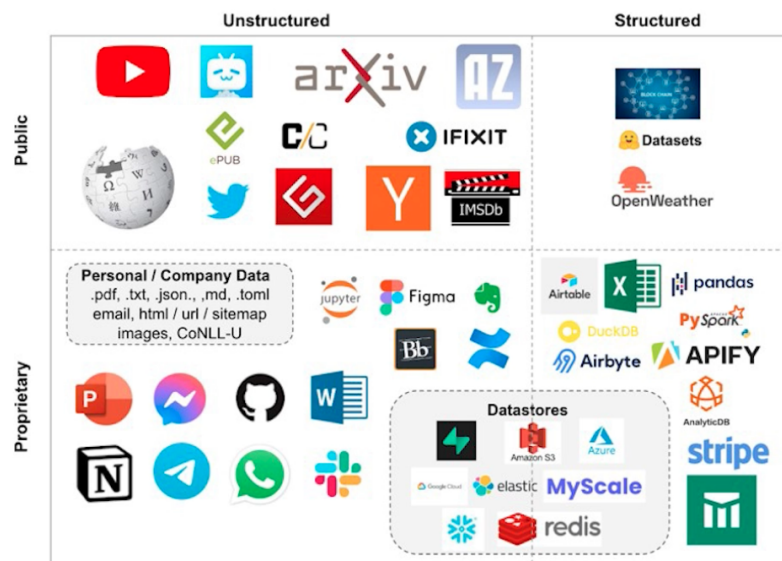
[Introduction](#)
[Document loading](#)
[Document splitting](#)
[Vector stores and embeddings](#)
[Retrieval](#)
[Question answering](#)

Introduction

Welcome 🤗

Document loading

- document loaders



- PDF
 - when loading a PDF, it will be split into pages
 - each page contains text, and metadata about the document it is a part of
- youtube video transcript
 - OpenAI Whisper
- URLs
 - `WebBaseLoad(f"{URL}")`
- Notion

- Notion directory loaders

Document splitting

- you need to split documents into smaller chunks
 - `chunk_size` & `chunk_overlap`
- types of splitters
 - `text_splitters`
- the metadata of sources and source chunks are the same
- context-aware splitting
 - you can split .md files based on headers and sub-headers

Vector stores and embeddings

- use of embeddings → vectors
- ChromaDB is in-memory and lightweight ⇒ perfect for prototyping
- edge cases and failures
 - duplicate documents
 - “third lecture” doesn’t quite work for embeddings ⇒ you get things that aren’t related to this one specifically

Retrieval

- query → retrieve more relevant chunk
- maximum marginal relevance (MMR)
 - sometimes we want to include information that is not the most similar to the query (eg toxicity of mushrooms)
 - get top fetch_K most relevant responses, then choose K responses that should be the most diverse
- self-query (LLM aided retrieval)
- compression: reduce size of output
- we can retrieve without a vector database (TF-IDF, SVM) ⇒ `tfidf retriever`, `svm retriever`

Question answering

- additional methods
 - `map_reduce`
 - `refine`
 - `map_rerank`